

The agent that asks about the *word*. Never the sentence.

A voice agent that catches the one word it misheard and asks a single, surgical question to fix it — instead of making the caller repeat everything.

You call your dentist. Traffic behind you, the line isn't great.
You say —

“Thursday.”

EVERY OTHER AGENT

“Sorry, I didn't catch that — could you repeat that?”

SAY LESS

“Tuesday or Thursday?”

You say one word back. You're booked. That's the whole idea.

THE INSIGHT

Confidence lies. Membership doesn't.

Speech models are often confidently wrong about a single hard word — a name, a day, a near-miss like “kid's cut” vs “kids cut.” No confidence score catches that. But checking whether the word it heard is even a valid answer does. If it isn't on the list, Say Less knows something's off — before it ever books the wrong thing.

Three moves, the instant your turn ends.

01

Catch

The transcribed value isn't on the field's allowed list — a day, a service, a surname — even though the model sounded sure. That mismatch is the tell.

02

Recover

Say Less searches the real menu of valid answers for the closest phonetic match — how the word sounds, not how it's spelled.

03

Ask small

It asks the most specific question the evidence supports, and nothing bigger — a single name, then a pair, then a category, only escalating if it must.

The specificity ladder — confident → unsure

“Kids cut?”

Name one · Say Less starts here

“Tuesday or Thursday?”

Offer two

“Which day works?”

Ask the group

“Sorry, could you repeat that?”

Ask again · everyone else lives here

THE PROOF

93 → 100%

Booking accuracy, same agent, same calls — the only thing that changes is turning the repair logic on.

0.07

words re-said per booking, vs 6–12 words on a typical repeat-the-sentence agent.

667

wrong bookings caught and fixed per 10,000 inbound calls — ones a normal agent would have silently shipped.

\$4.7k–\$8k

in avoided human clean-up cost per 10,000 calls, at typical contact-center recovery rates.

Measured on real recorded calls plus the public SLURP corpus, alongside a phone-degraded test set. Tested on synthetic degraded codec speech; confidently-wrong in-set tokens remain a residual risk, stated openly.

Built on the signal most agents throw away.

AssemblyAI Universal-Streaming

Real-time transcription with per-word confidence — the raw signal a normal agent discards after the transcript is final.

AssemblyAI LLM Gateway

Binds the live transcript onto bounded fields using paired clean-and-garbled few-shot examples.

Phonetic matching

Double Metaphone plus fuzzy matching finds the closest real answer on your own list — never a made-up guess.

Reactive keyterms

Narrowed candidates are pushed back into the live stream as a keyterms prompt before the retry audio even arrives.

89 / 89

unit tests passing

0

network calls to reproduce results

4.5s

to replay the full benchmark offline

Conversational repair is the *missing primitive* in voice AI.

When speech recognition meets a bounded, real-world workflow, intelligent repair is what bridges a raw transcript to a flawless action.